

# *dynoNet*: a neural network architecture for learning dynamical systems

Marco Forgione, Dario Piga

<sup>1</sup>IDSIA Dalle Molle Institute for Artificial Intelligence SUPSI-USI, Lugano, Switzerland

Convegno Nazionale Automatica  
Cagliari, 9–11 Settembre 2020

# Motivations

Two main classes of **neural network** structures for sequence modeling and system identification:

## Recurrent NNs

General state-space models

- High representational capacity
- Hard to parallelize
- Numerical issues in training

## 1D Convolutional NNs

Dynamics through FIR blocks

- Lower capacity, several params
- Fully parallelizable
- Fast, well-behaved training

We introduce *dynoNet*: an architecture using linear dynamical operators parametrized as **rational transfer functions** as building blocks.

- Extends 1D Convolutional NNs to **Infinite Impulse Response** dynamics
- Can be trained efficiently by plain **back-propagation**

# Motivations

Two main classes of **neural network** structures for sequence modeling and system identification:

## Recurrent NNs

General state-space models

- High representational capacity
- Hard to parallelize
- Numerical issues in training

## 1D Convolutional NNs

Dynamics through FIR blocks

- Lower capacity, several params
- Fully parallelizable
- Fast, well-behaved training

We introduce *dynoNet*: an architecture using linear dynamical operators parametrized as **rational transfer functions** as building blocks.

- Extends 1D Convolutional NNs to **Infinite Impulse Response** dynamics
- Can be trained efficiently by plain **back-propagation**

# Motivations

Two main classes of **neural network** structures for sequence modeling and system identification:

## Recurrent NNs

General state-space models

- High representational capacity
- Hard to parallelize
- Numerical issues in training

## 1D Convolutional NNs

Dynamics through FIR blocks






- Lower capacity, several params
- Fully parallelizable
- Fast, well-behaved training

We introduce *dynoNet*: an architecture using linear dynamical operators parametrized as **rational transfer functions** as building blocks.

- Extends 1D Convolutional NNs to **Infinite Impulse Response** dynamics
- Can be trained efficiently by plain **back-propagation**

## Related works

**Block-oriented** architectures consist in the interconnection of transfer functions  $G(z)$  and static non-linearities  $F(\cdot)$ :








extensively studied in System Identification.

Training with **specialized algorithms** requiring, e.g. analytic expressions of gradients/jacobians.

## Related works

**Block-oriented** architectures consist in the interconnection of transfer functions  $G(z)$  and static non-linearities  $F(\cdot)$ :




extensively studied in System Identification.

Training with **specialized algorithms** requiring, e.g. analytic expressions of gradients/jacobians.


# *dynoNet*

- *dynoNet* generalizes block-oriented models to **arbitrary connection** of MIMO blocks  $G(z)$  and  $F(\cdot)$
- More importantly, training is performed using a **general approach**
- Plain **back-propagation** for gradient computation exploiting Deep Learning software



**Technical challenge:** back-propagation through the transfer function!  
No hint in the literature, no ready-made implementation available.

- *dynoNet* generalizes block-oriented models to **arbitrary connection** of MIMO blocks  $G(z)$  and  $F(\cdot)$
- More importantly, training is performed using a **general approach**
- Plain **back-propagation** for gradient computation exploiting Deep Learning software



**Technical challenge:** back-propagation through the transfer function!  
No hint in the literature, no ready-made implementation available.



# Transfer function (SISO)

Transforms an input sequence  $u(t)$  to an output  $y(t)$  according to:

$$y(t) = G(q)u(t) = \frac{b_0 + b_1q^{-1} + \dots + b_{n_b}q^{-n_b}}{1 + a_1q^{-1} + \dots + a_{n_a}q^{-n_a}}u(t)$$

Equivalent to the recurrence equation:

$$y(t) = b_0u(t) + b_1u(t-1) + \dots + b_{n_b}u(t-n_b) - a_1y(t-1) \dots - a_{n_a}y(t-n_a).$$

For our purposes,  $G$  is a **vector operator** with coefficients  $a, b$ , transforming  $\mathbf{u} \in \mathbb{R}^T$  to  $\mathbf{y} \in \mathbb{R}^T$

$$\mathbf{y} = G(\mathbf{u}; a, b)$$

Our goal is to provide  $G$  with a **back-propagation** behavior.  
The operation has to be **efficient**!

# Transfer function (SISO)

Transforms an input sequence  $u(t)$  to an output  $y(t)$  according to:

$$y(t) = G(q)u(t) = \frac{b_0 + b_1q^{-1} + \dots + b_{n_b}q^{-n_b}}{1 + a_1q^{-1} + \dots + a_{n_a}q^{-n_a}}u(t)$$

Equivalent to the recurrence equation:

$$y(t) = b_0u(t) + b_1u(t-1) + \dots + b_{n_b}u(t-n_b) - a_1y(t-1) \dots - a_{n_a}y(t-n_a).$$

For our purposes,  $G$  is a **vector operator** with coefficients  $a, b$ , transforming  $\mathbf{u} \in \mathbb{R}^T$  to  $\mathbf{y} \in \mathbb{R}^T$

$$\mathbf{y} = G(\mathbf{u}; a, b)$$

Our goal is to provide  $G$  with a **back-propagation** behavior.  
The operation has to be **efficient**!

# Transfer function (SISO)

Transforms an input sequence  $u(t)$  to an output  $y(t)$  according to:

$$y(t) = G(q)u(t) = \frac{b_0 + b_1q^{-1} + \dots + b_{n_b}q^{-n_b}}{1 + a_1q^{-1} + \dots + a_{n_a}q^{-n_a}}u(t)$$

Equivalent to the recurrence equation:

$$y(t) = b_0u(t) + b_1u(t-1) + \dots + b_{n_b}u(t-n_b) - a_1y(t-1) \dots - a_{n_a}y(t-n_a).$$

For our purposes,  $G$  is a **vector operator** with coefficients  $a, b$ , transforming  $\mathbf{u} \in \mathbb{R}^T$  to  $\mathbf{y} \in \mathbb{R}^T$

$$\mathbf{y} = G(\mathbf{u}; a, b)$$

Our goal is to provide  $G$  with a **back-propagation** behavior.  
The operation has to be **efficient**!


# Forward pass

In back-propagation-based training, the user defines a **computational graph** producing a **loss**  $\mathcal{L}$  (to be minimized).

In the **forward pass**, the loss  $\mathcal{L}$  is computed.

$G$  receives  $\mathbf{u}$ ,  $a$ , and  $b$  and needs to compute  $\mathbf{y}$ :

$$\mathbf{y} = G.\text{forward}(\mathbf{u}; a, b).$$



The forward pass for  $G$  is easy: it is just the filtering operation!

Computational cost:  $\mathcal{O}(T)$ .


# Forward pass

In back-propagation-based training, the user defines a **computational graph** producing a **loss**  $\mathcal{L}$  (to be minimized).

In the **forward pass**, the loss  $\mathcal{L}$  is computed.

$G$  receives  $\mathbf{u}$ ,  $a$ , and  $b$  and needs to compute  $\mathbf{y}$ :

$$\mathbf{y} = G.\text{forward}(\mathbf{u}; a, b).$$



The forward pass for  $G$  is easy: it is just the filtering operation!

Computational cost:  $\mathcal{O}(T)$ .

# Backward pass

- In the **backward pass**, derivatives of  $\mathcal{L}$  w.r.t. the **training variables** are computed. Notation:  $\bar{x} = \frac{\partial \mathcal{L}}{\partial x}$ .
- The procedure starts from  $\bar{\mathcal{L}} \equiv \frac{\partial \mathcal{L}}{\partial \mathcal{L}} = 1$  and goes **backward**.
- Each operator must be able to “push back” derivatives from its outputs to its inputs

$G$  receives  $\bar{y} \equiv \frac{\partial \mathcal{L}}{\partial y}$  and is responsible for computing:  $\bar{u}, \bar{a}, \bar{b}$ :

$$\bar{u}, \bar{a}, \bar{b} = G.\text{backward}(\bar{y}; a, b).$$


Chain rule of calculus is the basic ingredient, but certain tricks may be used to **speed up** the operation. Let us see an example...

# Backward pass

- In the **backward pass**, derivatives of  $\mathcal{L}$  w.r.t. the **training variables** are computed. Notation:  $\bar{x} = \frac{\partial \mathcal{L}}{\partial x}$ .
- The procedure starts from  $\bar{\mathcal{L}} \equiv \frac{\partial \mathcal{L}}{\partial \mathcal{L}} = 1$  and goes **backward**.
- Each operator must be able to “push back” derivatives from its outputs to its inputs

$G$  receives  $\bar{\mathbf{y}} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{y}}$  and is responsible for computing:  $\bar{\mathbf{u}}, \bar{a}, \bar{b}$ :

$$\bar{\mathbf{u}}, \bar{a}, \bar{b} = G.\text{backward}(\bar{\mathbf{y}}; a, b).$$




Chain rule of calculus is the basic ingredient, but certain tricks may be used to **speed up** the operation. Let us see an example...

# Backward pass

- In the **backward pass**, derivatives of  $\mathcal{L}$  w.r.t. the **training variables** are computed. Notation:  $\bar{x} = \frac{\partial \mathcal{L}}{\partial x}$ .
- The procedure starts from  $\bar{\mathcal{L}} \equiv \frac{\partial \mathcal{L}}{\partial \mathcal{L}} = 1$  and goes **backward**.
- Each operator must be able to “push back” derivatives from its outputs to its inputs

$G$  receives  $\bar{\mathbf{y}} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{y}}$  and is responsible for computing:  $\bar{\mathbf{u}}, \bar{a}, \bar{b}$ :

$$\bar{\mathbf{u}}, \bar{a}, \bar{b} = G.\text{backward}(\bar{\mathbf{y}}; a, b).$$



**Chain rule** of calculus is the basic ingredient, but certain tricks may be used to **speed up** the operation. Let us see an example...



## Backward pass for $\mathbf{u}$

Compute  $\bar{\mathbf{u}} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{u}}$  from  $\bar{\mathbf{y}} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{y}}$ .

- Applying the chain rule:

$$\bar{\mathbf{u}}_\tau = \frac{\partial \mathcal{L}}{\partial \mathbf{u}_\tau} = \sum_{t=0}^{T-1} \frac{\partial \mathcal{L}}{\partial \mathbf{y}_t} \frac{\partial \mathbf{y}_t}{\partial \mathbf{u}_\tau} = \sum_{t=0}^{T-1} \bar{\mathbf{y}}_t \mathbf{g}_{t-\tau}$$

where  $\mathbf{g}$  is the **impulse response** of  $G$ .

- From the expression above, by definition:

$$\bar{\mathbf{u}} = \mathbf{g} \star \bar{\mathbf{y}},$$

where  $\star$  is **cross-correlation**. This implementation has cost  $\mathcal{O}(T^2)$

- It is equivalent to filtering  $\bar{\mathbf{y}}$  through  $G$  in reverse time, and flipping the result. Implemented this way, the cost is  $\mathcal{O}(T)$ !

$$\bar{\mathbf{u}} = \text{flip}(G(q)\text{flip}(\bar{\mathbf{y}}))$$

All details also for  $\bar{a}$  and  $\bar{b}$  in the *dynoNet* arXiv paper...

## Backward pass for $\mathbf{u}$

Compute  $\bar{\mathbf{u}} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{u}}$  from  $\bar{\mathbf{y}} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{y}}$ .

- Applying the chain rule:

$$\bar{\mathbf{u}}_\tau = \frac{\partial \mathcal{L}}{\partial \mathbf{u}_\tau} = \sum_{t=0}^{T-1} \frac{\partial \mathcal{L}}{\partial \mathbf{y}_t} \frac{\partial \mathbf{y}_t}{\partial \mathbf{u}_\tau} = \sum_{t=0}^{T-1} \bar{\mathbf{y}}_t \mathbf{g}_{t-\tau}$$

where  $\mathbf{g}$  is the **impulse response** of  $G$ .

- From the expression above, by definition:

$$\bar{\mathbf{u}} = \mathbf{g} \star \bar{\mathbf{y}},$$

where  $\star$  is **cross-correlation**. This implementation has cost  $\mathcal{O}(T^2)$

- It is equivalent to filtering  $\bar{\mathbf{y}}$  through  $G$  in reverse time, and flipping the result. Implemented this way, the cost is  $\mathcal{O}(T)$ !

$$\bar{\mathbf{u}} = \text{flip}(G(q)\text{flip}(\bar{\mathbf{y}}))$$

All details also for  $\bar{a}$  and  $\bar{b}$  in the *dynoNet* arXiv paper...

## Backward pass for $\mathbf{u}$

Compute  $\bar{\mathbf{u}} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{u}}$  from  $\bar{\mathbf{y}} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{y}}$ .

- Applying the chain rule:

$$\bar{\mathbf{u}}_\tau = \frac{\partial \mathcal{L}}{\partial \mathbf{u}_\tau} = \sum_{t=0}^{T-1} \frac{\partial \mathcal{L}}{\partial \mathbf{y}_t} \frac{\partial \mathbf{y}_t}{\partial \mathbf{u}_\tau} = \sum_{t=0}^{T-1} \bar{\mathbf{y}}_t \mathbf{g}_{t-\tau}$$

where  $\mathbf{g}$  is the **impulse response** of  $G$ .

- From the expression above, by definition:

$$\bar{\mathbf{u}} = \mathbf{g} \star \bar{\mathbf{y}},$$

where  $\star$  is **cross-correlation**. This implementation has cost  $\mathcal{O}(T^2)$

- It is equivalent to filtering  $\bar{\mathbf{y}}$  through  $G$  in reverse time, and flipping the result. Implemented this way, the cost is  $\mathcal{O}(T)$ !

$$\bar{\mathbf{u}} = \text{flip}(G(q)\text{flip}(\bar{\mathbf{y}}))$$

All details also for  $\bar{a}$  and  $\bar{b}$  in the *dynoNet* arXiv paper...

## Backward pass for $\mathbf{u}$

Compute  $\bar{\mathbf{u}} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{u}}$  from  $\bar{\mathbf{y}} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{y}}$ .

- Applying the chain rule:

$$\bar{\mathbf{u}}_\tau = \frac{\partial \mathcal{L}}{\partial \mathbf{u}_\tau} = \sum_{t=0}^{T-1} \frac{\partial \mathcal{L}}{\partial \mathbf{y}_t} \frac{\partial \mathbf{y}_t}{\partial \mathbf{u}_\tau} = \sum_{t=0}^{T-1} \bar{\mathbf{y}}_t \mathbf{g}_{t-\tau}$$

where  $\mathbf{g}$  is the **impulse response** of  $G$ .

- From the expression above, by definition:

$$\bar{\mathbf{u}} = \mathbf{g} \star \bar{\mathbf{y}},$$

where  $\star$  is **cross-correlation**. This implementation has cost  $\mathcal{O}(T^2)$

- It is equivalent to filtering  $\bar{\mathbf{y}}$  through  $G$  in reverse time, and flipping the result. Implemented this way, the cost is  $\mathcal{O}(T)$ !

$$\bar{\mathbf{u}} = \text{flip}(G(q)\text{flip}(\bar{\mathbf{y}}))$$

All details also for  $\bar{a}$  and  $\bar{b}$  in the *dynoNet* arXiv paper...

## Backward pass for $\mathbf{u}$

Compute  $\bar{\mathbf{u}} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{u}}$  from  $\bar{\mathbf{y}} \equiv \frac{\partial \mathcal{L}}{\partial \mathbf{y}}$ .

- Applying the chain rule:

$$\bar{\mathbf{u}}_\tau = \frac{\partial \mathcal{L}}{\partial \mathbf{u}_\tau} = \sum_{t=0}^{T-1} \frac{\partial \mathcal{L}}{\partial \mathbf{y}_t} \frac{\partial \mathbf{y}_t}{\partial \mathbf{u}_\tau} = \sum_{t=0}^{T-1} \bar{\mathbf{y}}_t \mathbf{g}_{t-\tau}$$

where  $\mathbf{g}$  is the **impulse response** of  $G$ .

- From the expression above, by definition:

$$\bar{\mathbf{u}} = \mathbf{g} \star \bar{\mathbf{y}},$$

where  $\star$  is **cross-correlation**. This implementation has cost  $\mathcal{O}(T^2)$

- It is equivalent to filtering  $\bar{\mathbf{y}}$  through  $G$  in reverse time, and flipping the result. Implemented this way, the cost is  $\mathcal{O}(T)$ !

$$\bar{\mathbf{u}} = \text{flip}(G(q)\text{flip}(\bar{\mathbf{y}}))$$

All details also for  $\bar{a}$  and  $\bar{b}$  in the *dynoNet* arXiv paper...


# PyTorch implementation

PyTorch implementation of the  $G$ -block in the repository

<https://github.com/forgi86/dynonet>.

Use case:

*dynoNet* architecture



Python code

```
G1 = LinearMimo(1, 4, ...) # a SIMO tf
F = StaticNonLin(4, 3, ...) # a static NN
G2 = LinearMimo(3, 1, ...) # a MISO tf
G3 = LinearMimo(1, 1, ...) # a SISO tf
```

```
def model(in_data):
    y1 = G1(in_data)
    z1 = F(y1)
    y2 = G2(z1)
    out = y2 + G3(in_data)
```

Any gradient-based optimization algorithm can be used to train the *dynoNet* with derivatives readily obtained by back-propagation.


# PyTorch implementation

PyTorch implementation of the  $G$ -block in the repository

<https://github.com/forgi86/dynonet>.

Use case:

*dynoNet* architecture



Python code

```
G1 = LinearMimo(1, 4, ...) # a SIMO tf
F = StaticNonLin(4, 3, ...) # a static NN
G2 = LinearMimo(3, 1, ...) # a MISO tf
G3 = LinearMimo(1, 1, ...) # a SISO tf
```


```
def model(in_data):
    y1 = G1(in_data)
    z1 = F(y1)
    y2 = G2(z1)
    out = y2 + G3(in_data)
```

Any **gradient-based** optimization algorithm can be used to train the *dynoNet* with derivatives readily obtained by **back-propagation**.

# Experimental results


Numerical experiments on public system identification benchmark available at [www.nonlinearbenchmark.org](http://www.nonlinearbenchmark.org).

Wiener-Hammerstein




fit = 99.5%

Bouc-Wen



fit = 93.2%

EMPS



fit = 96.8%


Compare favorably with state-of-the-art black-box identification techniques.



# Experimental results


Numerical experiments on public system identification benchmark available at [www.nonlinearbenchmark.org](http://www.nonlinearbenchmark.org).

Wiener-Hammerstein




fit = 99.5%

Bouc-Wen



fit = 93.2%

EMPS



fit = 96.8%

Compare favorably with state-of-the-art black-box identification techniques.

# Conclusions

A neural network architecture containing linear dynamical operators parametrized as rational transfer functions.

- Extends **1D-Convolutional** NNs to Infinite Impulse Response dynamics
- Extends **block-oriented** dynamical models with arbitrary interconnections
- Enables training through **plain back-propagation**, at cost  $\mathcal{O}(T)$ .  
No custom algorithm/code required

Current and future work:

- Estimation/control strategies
- System analysis/model reduction using e.g. linear tools

# Conclusions

A neural network architecture containing linear dynamical operators parametrized as rational transfer functions.

- Extends **1D-Convolutional** NNs to Infinite Impulse Response dynamics
- Extends **block-oriented** dynamical models with arbitrary interconnections
- Enables training through **plain back-propagation**, at cost  $\mathcal{O}(T)$ .  
No custom algorithm/code required

Current and future work:

- Estimation/control strategies
- System analysis/model reduction using e.g. linear tools

Thank you.  
Questions?

`marco.forgione@idsia.ch`